

Victor Chua

Singapore | +65 9190 5651 | chua1102@e.ntu.edu.sg | github.com/zeppyclown

Jason Leen | FinProjections, a company by Udiyo Pte Ltd

12 August 2026

Dear Jason,

You asked for a backend or data project where the database design was mine, so here is that first.

I built and own the Postgres layer behind Hermes, a dashboard that ingests career emails and job listings and serves the application documents I generate from them: github.com/ZeppyClown/email_dashboard. Read `schema.sql` rather than this letter. Three decisions in it are the ones I would defend. The primary key is a content hash of the source listing, not a random UUID, because the scraper re-runs over overlapping windows and the same job has to land on the same row instead of quietly duplicating. Each record is split on purpose: fields I filter or constrain are typed columns with CHECK constraints and indexes, while per-category free-form detail sits in one JSONB column, so adding a fact costs no migration and no table-per-category sprawl. Access is enforced in the database, not the client: row-level security allows public read, but updates are gated on a specific JWT email claim rather than `auth.role()`, because any stranger who requests a magic link for their own address would satisfy a role check. Inserts and deletes are revoked from browser roles, and the update grant is column-scoped to two columns out of sixteen.

That is the closest I have come to what you describe, and I should be straight that it is single-owner. I have not designed tenant isolation at scale, only the mechanism underneath it.

On accuracy: reviewing an evaluation pipeline at my current internship, I found the headline figure had been measured on the same examples used to ground the judge, so it was not a blind test. I re-reported it on held-out data and labelled the old number, even though the better-looking one was already in circulation. I also built their invoice pipeline, extracting vendor, amount and billing period from PDFs on-premises so financial documents never leave the network.

I read and write plain JavaScript with no framework, and I write schemas as their own data dictionary, each column carrying why it exists and what I rejected. I use MCP daily as a consumer but have not built a server; structuring data so agents can use it reliably is the part of this role I most want to learn.

One honest flag: the listing says ASAP and full time, and I return to NTU for my third year this month. I can commit consistent remote hours on a schedule agreed up front, but I would rather say that now than disappoint you in week two. If the hours are flexible, I would like to talk.

Yours sincerely,
Victor Chua